

HAR_LSTM ASSIGNMENT

September 25, 2018

```
In [1]: # Importing Libraries

In [2]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')

In [3]: # Activities are the class labels
# It is a 6 class classification
ACTIVITIES = {
    0: 'WALKING',
    1: 'WALKING_UPSTAIRS',
    2: 'WALKING_DOWNSTAIRS',
    3: 'SITTING',
    4: 'STANDING',
    5: 'LAYING',
}
```

0.0.1 Data

```
In [4]: # Data directory
DATADIR = 'UCI_HAR_Dataset'

In [5]: # Raw data signals
# Signals are from Accelerometer and Gyroscope
# The signals are in x,y,z directions
# Sensor signals are filtered to have only body acceleration
# excluding the acceleration due to gravity
# Triaxial acceleration from the accelerometer is total acceleration
SIGNALS = [
    "body_acc_x",
    "body_acc_y",
    "body_acc_z",
    "body_gyro_x",
    "body_gyro_y",
    "body_gyro_z",
    "total_acc_x",
    "total_acc_y",
```

```

        "total_acc_z"
    ]

In [6]: # Utility function to read the data from csv file
def _read_csv(filename):
    return pd.read_csv(filename, delim_whitespace=True, header=None)

# Utility function to load the load
def load_signals(subset):
    signals_data = []

    for signal in SIGNALS:
        filename = f'UCI_HAR_Dataset/{subset}/Inertial Signals/{signal}_{subset}.txt'
        signals_data.append(
            _read_csv(filename).as_matrix()
        )

    # Transpose is used to change the dimensionality of the output,
    # aggregating the signals by combination of sample/timestep.
    # Resultant shape is (7352 train/2947 test samples, 128 timesteps, 9 signals)
    return np.transpose(signals_data, (1, 2, 0))

In [7]: def load_y(subset):
    """
    The objective that we are trying to predict is a integer, from 1 to 6,
    that represents a human activity. We return a binary representation of
    every sample objective as a 6 bits vector using One Hot Encoding
    (https://pandas.pydata.org/pandas-docs/stable/generated/pandas.get\_dummies.html)
    """
    filename = f'UCI_HAR_Dataset/{subset}/y_{subset}.txt'
    y = _read_csv(filename)[0]

    return pd.get_dummies(y).as_matrix()

In [8]: def load_data():
    """
    Obtain the dataset from multiple files.
    Returns: X_train, X_test, y_train, y_test
    """
    X_train, X_test = load_signals('train'), load_signals('test')
    y_train, y_test = load_y('train'), load_y('test')

    return X_train, X_test, y_train, y_test

In [9]: # Importing tensorflow
np.random.seed(42)
import tensorflow as tf
tf.set_random_seed(42)

```

```

In [10]: # Configuring a session
        session_conf = tf.ConfigProto(
            intra_op_parallelism_threads=1,
            inter_op_parallelism_threads=1
        )

In [11]: # Import Keras
        from keras import backend as K
        sess = tf.Session(graph=tf.get_default_graph(), config=session_conf)
        K.set_session(sess)

```

Using TensorFlow backend.

```

In [12]: # Importing libraries
        from keras.models import Sequential
        from keras.layers import LSTM
        from keras.layers.core import Dense, Dropout

In [13]: # Initializing parameters
        epochs = 30
        batch_size = 16
        n_hidden = 32

In [14]: # Utility function to count the number of classes
        def _count_classes(y):
            return len(set([tuple(category) for category in y]))

In [15]: # Loading the train and test data
        X_train, X_test, Y_train, Y_test = load_data()

In [16]: timesteps = len(X_train[0])
        input_dim = len(X_train[0][0])
        n_classes = _count_classes(Y_train)

        print(timesteps)
        print(input_dim)
        print(len(X_train))

```

```

128
9
7352

```

- Defining the Architecture of LSTM

1 (1) Model having 1 LSTM layer with 32 LSTM Units

```

In [17]: # Initiliazing the sequential model
        model = Sequential()

```

```

# Configuring the parameters
model.add(LSTM(n_hidden, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model.add(Dense(n_classes, activation='sigmoid'))
model.summary()

# Compiling the model
model.compile(loss='categorical_crossentropy',optimizer='rmsprop',metrics=['accuracy'])

# Training the model
history = model.fit(X_train, Y_train, batch_size=batch_size,validation_data=(X_test, '

```

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 32)	5376
dropout_1 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 6)	198

Total params: 5,574
 Trainable params: 5,574
 Non-trainable params: 0

Train on 7352 samples, validate on 2947 samples

Epoch	Train Samples	Train Loss	Train Acc	Val Samples	Val Loss	Val Acc
1/30	7352/7352	1.3139	0.4358	2947	1.3139	0.4358
2/30	7352/7352	0.9788	0.5773	2947	0.9788	0.5773
3/30	7352/7352	0.7977	0.6457	2947	0.7977	0.6457
4/30	7352/7352	0.6989	0.6582	2947	0.6989	0.6582
5/30	7352/7352	0.6359	0.6797	2947	0.6359	0.6797
6/30	7352/7352	0.5819	0.6865	2947	0.5819	0.6865
7/30	7352/7352	0.5676	0.7058	2947	0.5676	0.7058
8/30	7352/7352	0.5583	0.7217	2947	0.5583	0.7217
9/30	7352/7352	0.5386	0.7557	2947	0.5386	0.7557
10/30	7352/7352	0.4804	0.7911	2947	0.4804	0.7911

```

Epoch 11/30
7352/7352 [=====] - 37s 5ms/step - loss: 0.4320 - acc: 0.8052 - val_loss: 0.4320
Epoch 12/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.4279 - acc: 0.8062 - val_loss: 0.4279
Epoch 13/30
7352/7352 [=====] - 38s 5ms/step - loss: 0.3911 - acc: 0.8130 - val_loss: 0.3911
Epoch 14/30
7352/7352 [=====] - 37s 5ms/step - loss: 0.3898 - acc: 0.8313 - val_loss: 0.3898
Epoch 15/30
7352/7352 [=====] - 37s 5ms/step - loss: 0.3308 - acc: 0.8942 - val_loss: 0.3308
Epoch 16/30
7352/7352 [=====] - 37s 5ms/step - loss: 0.2891 - acc: 0.9176 - val_loss: 0.2891
Epoch 17/30
7352/7352 [=====] - 38s 5ms/step - loss: 0.2660 - acc: 0.9246 - val_loss: 0.2660
Epoch 18/30
7352/7352 [=====] - 38s 5ms/step - loss: 0.2538 - acc: 0.9251 - val_loss: 0.2538
Epoch 19/30
7352/7352 [=====] - 38s 5ms/step - loss: 0.2502 - acc: 0.9312 - val_loss: 0.2502
Epoch 20/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.1980 - acc: 0.9382 - val_loss: 0.1980
Epoch 21/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2018 - acc: 0.9372 - val_loss: 0.2018
Epoch 22/30
7352/7352 [=====] - 39s 5ms/step - loss: 0.2455 - acc: 0.9310 - val_loss: 0.2455
Epoch 23/30
7352/7352 [=====] - 40s 5ms/step - loss: 0.2194 - acc: 0.9329 - val_loss: 0.2194
Epoch 24/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2282 - acc: 0.9304 - val_loss: 0.2282
Epoch 25/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.2166 - acc: 0.9359 - val_loss: 0.2166
Epoch 26/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2173 - acc: 0.9350 - val_loss: 0.2173
Epoch 27/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.2224 - acc: 0.9353 - val_loss: 0.2224
Epoch 28/30
7352/7352 [=====] - 39s 5ms/step - loss: 0.1961 - acc: 0.9385 - val_loss: 0.1961
Epoch 29/30
7352/7352 [=====] - 39s 5ms/step - loss: 0.1876 - acc: 0.9416 - val_loss: 0.1876
Epoch 30/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1999 - acc: 0.9411 - val_loss: 0.1999

```

```

In [18]: import matplotlib.pyplot as plt
         %matplotlib inline
         import seaborn as sns
         from sklearn.metrics import confusion_matrix

         # Final evaluation of the model

```

```

scores = model.evaluate(X_test, Y_test, verbose=0)
print("Test Score: %f" % (scores[0]))
print("Test Accuracy: %f%%" % (scores[1]*100))

# Confusion Matrix
Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])
Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model.predict(X_test), axis=1)])


# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fontsize=12)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right', fontsize=12)
plt.ylabel('True label',size=18)
plt.xlabel('Predicted label',size=18)
plt.title("Confusion Matrix\n",size=24)
plt.show()

```

Test Score: 0.488270

Test Accuracy: 88.293180%



- With a simple 2 layer architecture we got 88.29% accuracy and a loss of 0.488
- We can further improve the performance with Hyperparameter tuning

2 (2) Model having 1 LSTM layer with 48 LSTM Units and 'adam' as an optimizer

```
In [19]: # Initiating the sequential model
model1 = Sequential()
# Configuring the parameters
model1.add(LSTM(48, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model1.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model1.add(Dense(n_classes, activation='sigmoid'))
print(model1.summary())
```

```

# Compiling the model
model1.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'])

# Training the model
history1 = model1.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 48)	11136
dropout_2 (Dropout)	(None, 48)	0
dense_2 (Dense)	(None, 6)	294

=====
 Total params: 11,430
 Trainable params: 11,430
 Non-trainable params: 0
 =====

None

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 45s 6ms/step - loss: 1.4210 - acc: 0.3677 - val_loss: 1.4210

Epoch 2/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.3615 - acc: 0.3659 - val_loss: 1.3615

Epoch 3/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.2965 - acc: 0.4147 - val_loss: 1.2965

Epoch 4/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.2413 - acc: 0.4645 - val_loss: 1.2413

Epoch 5/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.1199 - acc: 0.5102 - val_loss: 1.1199

Epoch 6/30

7352/7352 [=====] - 42s 6ms/step - loss: 1.0028 - acc: 0.5439 - val_loss: 1.0028

Epoch 7/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.0453 - acc: 0.5098 - val_loss: 1.0453

Epoch 8/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.1810 - acc: 0.4523 - val_loss: 1.1810

Epoch 9/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.2428 - acc: 0.4329 - val_loss: 1.2428

Epoch 10/30

7352/7352 [=====] - 43s 6ms/step - loss: 0.9496 - acc: 0.5747 - val_loss: 0.9496

Epoch 11/30

7352/7352 [=====] - 43s 6ms/step - loss: 1.0623 - acc: 0.5399 - val_loss: 1.0623

Epoch 12/30

7352/7352 [=====] - 43s 6ms/step - loss: 0.8686 - acc: 0.6114 - val_loss: 0.8686

Epoch 13/30

7352/7352 [=====] - 42s 6ms/step - loss: 1.0787 - acc: 0.4974 - val_loss: 1.0787

Epoch 14/30


```

7352/7352 [=====] - 43s 6ms/step - loss: 0.9513 - acc: 0.5822 - val_loss: 0.9513
Epoch 15/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.8773 - acc: 0.5929 - val_loss: 0.8773
Epoch 16/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.7541 - acc: 0.6250 - val_loss: 0.7541
Epoch 17/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7139 - acc: 0.6499 - val_loss: 0.7139
Epoch 18/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7097 - acc: 0.6468 - val_loss: 0.7097
Epoch 19/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.6794 - acc: 0.6575 - val_loss: 0.6794
Epoch 20/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.6810 - acc: 0.6553 - val_loss: 0.6810
Epoch 21/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.6905 - acc: 0.6468 - val_loss: 0.6905
Epoch 22/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.6648 - acc: 0.6712 - val_loss: 0.6648
Epoch 23/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.6891 - acc: 0.6727 - val_loss: 0.6891
Epoch 24/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.6327 - acc: 0.7331 - val_loss: 0.6327
Epoch 25/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.5341 - acc: 0.8074 - val_loss: 0.5341
Epoch 26/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.3767 - acc: 0.8696 - val_loss: 0.3767
Epoch 27/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.3015 - acc: 0.9015 - val_loss: 0.3015
Epoch 28/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.3263 - acc: 0.8989 - val_loss: 0.3263
Epoch 29/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.3337 - acc: 0.8976 - val_loss: 0.3337
Epoch 30/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.2294 - acc: 0.9272 - val_loss: 0.2294

```

```

In [20]: # Final evaluation of the model
scores1 = model1.evaluate(X_test, Y_test, verbose=0)
print("Test Score: %f" % (scores1[0]))
print("Test Accuracy: %f%%" % (scores1[1]*100))

# Confusion Matrix
Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])
Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model1.predict(X_test), axis=1)])

# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))

```

```
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```

```
# Setting tick labels for heatmap
```

```
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right',
```

```
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right',
```

```
plt.ylabel('True label',size=18)
```


```
plt.xlabel('Predicted label',size=18)
```

```
plt.title("Confusion Matrix\n",size=24)
```

```
plt.show()
```

Test Score: 0.344224

Test Accuracy: 87.139464%



3 (3) Model having 1 LSTM layer with 48 LSTM Units and 'rmsprop' as an optimizer

```
In [21]: # Initiliazing the sequential model
model2 = Sequential()
# Configuring the parameters
model2.add(LSTM(48, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model2.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model2.add(Dense(n_classes, activation='sigmoid'))
print(model2.summary())

# Compiling the model
model2.compile(loss='categorical_crossentropy',optimizer='rmsprop',metrics=['accuracy'])

# Training the model
history2 = model2.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_test))
```

```
-----
Layer (type)                 Output Shape              Param #
-----
lstm_3 (LSTM)                 (None, 48)                11136
-----
dropout_3 (Dropout)           (None, 48)                 0
-----
dense_3 (Dense)               (None, 6)                  294
=====
Total params: 11,430
Trainable params: 11,430
Non-trainable params: 0
-----
None
Train on 7352 samples, validate on 2947 samples
Epoch 1/30
7352/7352 [=====] - 43s 6ms/step - loss: 1.2313 - acc: 0.4780 - val_loss: 1.2313
Epoch 2/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.8782 - acc: 0.6073 - val_loss: 0.8782
Epoch 3/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.7840 - acc: 0.6542 - val_loss: 0.7840
Epoch 4/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.6928 - acc: 0.6900 - val_loss: 0.6928
Epoch 5/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.6225 - acc: 0.7348 - val_loss: 0.6225
Epoch 6/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.5056 - acc: 0.8290 - val_loss: 0.5056
Epoch 7/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.3532 - acc: 0.8900 - val_loss: 0.3532
```

Epoch 8/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.2994 - acc: 0.9113 - val_loss: 0.2994
Epoch 9/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2638 - acc: 0.9212 - val_loss: 0.2638
Epoch 10/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.2297 - acc: 0.9276 - val_loss: 0.2297
Epoch 11/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2190 - acc: 0.9336 - val_loss: 0.2190
Epoch 12/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2153 - acc: 0.9329 - val_loss: 0.2153
Epoch 13/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2055 - acc: 0.9376 - val_loss: 0.2055
Epoch 14/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1898 - acc: 0.9366 - val_loss: 0.1898
Epoch 15/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.2032 - acc: 0.9319 - val_loss: 0.2032
Epoch 16/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1801 - acc: 0.9416 - val_loss: 0.1801
Epoch 17/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1810 - acc: 0.9423 - val_loss: 0.1810
Epoch 18/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1714 - acc: 0.9452 - val_loss: 0.1714
Epoch 19/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1654 - acc: 0.9411 - val_loss: 0.1654
Epoch 20/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1795 - acc: 0.9455 - val_loss: 0.1795
Epoch 21/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1676 - acc: 0.9404 - val_loss: 0.1676
Epoch 22/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1811 - acc: 0.9423 - val_loss: 0.1811
Epoch 23/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1563 - acc: 0.9449 - val_loss: 0.1563
Epoch 24/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1495 - acc: 0.9449 - val_loss: 0.1495
Epoch 25/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1740 - acc: 0.9436 - val_loss: 0.1740
Epoch 26/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1564 - acc: 0.9446 - val_loss: 0.1564
Epoch 27/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1648 - acc: 0.9475 - val_loss: 0.1648
Epoch 28/30
7352/7352 [=====] - 42s 6ms/step - loss: 0.1504 - acc: 0.9438 - val_loss: 0.1504
Epoch 29/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1501 - acc: 0.9468 - val_loss: 0.1501
Epoch 30/30
7352/7352 [=====] - 41s 6ms/step - loss: 0.1647 - acc: 0.9471 - val_loss: 0.1647

```

In [22]: # Final evaluation of the model
scores2 = model2.evaluate(X_test, Y_test, verbose=0)
print("Test Score: %f" % (scores2[0]))
print("Test Accuracy: %f%%" % (scores2[1]*100))

# Confusion Matrix
Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])
Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model2.predict(X_test), axis=1)])


# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fontsize=12)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right', fontsize=12)
plt.ylabel('True label', size=18)
plt.xlabel('Predicted label', size=18)
plt.title("Confusion Matrix\n", size=24)
plt.show()

```

Test Score: 0.410484

Test Accuracy: 90.498812%



4 (4) Model having 1 LSTM layer with 64 LSTM Units and 'rmsprop' as an optimizer

```
In [23]: # Initiliazing the sequential model
model3 = Sequential()
# Configuring the parameters
model3.add(LSTM(64, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model3.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model3.add(Dense(n_classes, activation='sigmoid'))
print(model3.summary())

# Compiling the model
model3.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

```
# Training the model
```

```
history3 = model3.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_test))
```

```
-----
Layer (type)                 Output Shape              Param #
=====
lstm_4 (LSTM)                (None, 64)                18944
-----
dropout_4 (Dropout)          (None, 64)                0
-----
dense_4 (Dense)              (None, 6)                 390
=====

Total params: 19,334
Trainable params: 19,334
Non-trainable params: 0

-----
None
Train on 7352 samples, validate on 2947 samples
Epoch 1/30
7352/7352 [=====] - 47s 6ms/step - loss: 1.2746 - acc: 0.4457 - val_loss: 1.2746
Epoch 2/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.9587 - acc: 0.6020 - val_loss: 0.9587
Epoch 3/30
7352/7352 [=====] - 46s 6ms/step - loss: 1.0225 - acc: 0.5890 - val_loss: 1.0225
Epoch 4/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.7561 - acc: 0.6812 - val_loss: 0.7561
Epoch 5/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.6203 - acc: 0.7402 - val_loss: 0.6203
Epoch 6/30
7352/7352 [=====] - 58s 8ms/step - loss: 0.4874 - acc: 0.8249 - val_loss: 0.4874
Epoch 7/30
7352/7352 [=====] - 48s 7ms/step - loss: 0.3588 - acc: 0.8905 - val_loss: 0.3588
Epoch 8/30
7352/7352 [=====] - 52s 7ms/step - loss: 0.2826 - acc: 0.9042 - val_loss: 0.2826
Epoch 9/30
7352/7352 [=====] - 49s 7ms/step - loss: 0.2855 - acc: 0.9033 - val_loss: 0.2855
Epoch 10/30
7352/7352 [=====] - 48s 7ms/step - loss: 0.2367 - acc: 0.9197 - val_loss: 0.2367
Epoch 11/30
7352/7352 [=====] - 48s 6ms/step - loss: 0.2891 - acc: 0.9064 - val_loss: 0.2891
Epoch 12/30
7352/7352 [=====] - 48s 7ms/step - loss: 0.2101 - acc: 0.9327 - val_loss: 0.2101
Epoch 13/30
7352/7352 [=====] - 49s 7ms/step - loss: 0.1883 - acc: 0.9309 - val_loss: 0.1883
Epoch 14/30
7352/7352 [=====] - 48s 6ms/step - loss: 0.1781 - acc: 0.9354 - val_loss: 0.1781
Epoch 15/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1812 - acc: 0.9344 - val_loss: 0.1812
```

```

Epoch 16/30
7352/7352 [=====] - 48s 6ms/step - loss: 0.1701 - acc: 0.9414 - val_loss: 0.1701
Epoch 17/30
7352/7352 [=====] - 48s 6ms/step - loss: 0.1603 - acc: 0.9446 - val_loss: 0.1603
Epoch 18/30
7352/7352 [=====] - 48s 7ms/step - loss: 0.1494 - acc: 0.9460 - val_loss: 0.1494
Epoch 19/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1555 - acc: 0.9445 - val_loss: 0.1555
Epoch 20/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1413 - acc: 0.9498 - val_loss: 0.1413
Epoch 21/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1674 - acc: 0.9444 - val_loss: 0.1674
Epoch 22/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1550 - acc: 0.9430 - val_loss: 0.1550
Epoch 23/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1551 - acc: 0.9450 - val_loss: 0.1551
Epoch 24/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.1679 - acc: 0.9440 - val_loss: 0.1679
Epoch 25/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1543 - acc: 0.9472 - val_loss: 0.1543
Epoch 26/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1457 - acc: 0.9459 - val_loss: 0.1457
Epoch 27/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1383 - acc: 0.9476 - val_loss: 0.1383
Epoch 28/30
7352/7352 [=====] - 46s 6ms/step - loss: 0.1412 - acc: 0.9508 - val_loss: 0.1412
Epoch 29/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1496 - acc: 0.9464 - val_loss: 0.1496
Epoch 30/30
7352/7352 [=====] - 47s 6ms/step - loss: 0.1439 - acc: 0.9490 - val_loss: 0.1439

```

```

In [24]: # Final evaluation of the model
scores3 = model3.evaluate(X_test, Y_test, verbose=0)
print("Test Score: %f" % (scores3[0]))
print("Test Accuracy: %f%%" % (scores3[1]*100))

# Confusion Matrix
Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])
Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model3.predict(X_test), axis=1)])

# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap

```




```

heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right',
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right',
plt.ylabel('True label',size=18)
plt.xlabel('Predicted label',size=18)
plt.title("Confusion Matrix\n",size=24)
plt.show()

```

Test Score: 0.299268

Test Accuracy: 92.059722%



5 (5) Model having 2 LSTM layer with 32 LSTM Units and 'rmsprop' as an optimizer

```

In [25]: # Initiliazing the sequential model
         model4 = Sequential()

```

```

# Configuring the parameters
model4.add(LSTM(32,return_sequences=True, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model4.add(Dropout(0.5))

# Configuring the parameters
model4.add(LSTM(32))
# Adding a dropout layer
model4.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model4.add(Dense(n_classes, activation='sigmoid'))
print(model4.summary())

# Compiling the model
model4.compile(loss='categorical_crossentropy',optimizer='rmsprop',metrics=['accuracy'])

# Training the model
history4 = model4.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_test))

```

```

-----
Layer (type)                 Output Shape              Param #
=====
lstm_5 (LSTM)                 (None, 128, 32)          5376
-----
dropout_5 (Dropout)           (None, 128, 32)          0
-----
lstm_6 (LSTM)                 (None, 32)                8320
-----
dropout_6 (Dropout)           (None, 32)                0
-----
dense_5 (Dense)               (None, 6)                 198
=====
Total params: 13,894
Trainable params: 13,894
Non-trainable params: 0
-----
None
Train on 7352 samples, validate on 2947 samples
Epoch 1/30
7352/7352 [=====] - 86s 12ms/step - loss: 1.2107 - acc: 0.5061 - val_loss: 1.2107
Epoch 2/30
7352/7352 [=====] - 84s 11ms/step - loss: 0.7991 - acc: 0.6766 - val_loss: 0.7991
Epoch 3/30
7352/7352 [=====] - 83s 11ms/step - loss: 0.6209 - acc: 0.7542 - val_loss: 0.6209
Epoch 4/30
7352/7352 [=====] - 83s 11ms/step - loss: 0.4968 - acc: 0.7802 - val_loss: 0.4968
Epoch 5/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.4323 - acc: 0.8009 - val_loss: 0.4323

```

Epoch 6/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.4538 - acc: 0.8247 - val_
Epoch 7/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.3524 - acc: 0.8785 - val_
Epoch 8/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.3199 - acc: 0.9115 - val_
Epoch 9/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.2609 - acc: 0.9285 - val_
Epoch 10/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.2290 - acc: 0.9339 - val_
Epoch 11/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.2160 - acc: 0.9353 - val_
Epoch 12/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.2236 - acc: 0.9321 - val_
Epoch 13/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1718 - acc: 0.9455 - val_
Epoch 14/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1740 - acc: 0.9359 - val_
Epoch 15/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1693 - acc: 0.9423 - val_
Epoch 16/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1813 - acc: 0.9459 - val_
Epoch 17/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1687 - acc: 0.9472 - val_
Epoch 18/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1557 - acc: 0.9474 - val_
Epoch 19/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1479 - acc: 0.9471 - val_
Epoch 20/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1479 - acc: 0.9493 - val_
Epoch 21/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1465 - acc: 0.9509 - val_
Epoch 22/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1508 - acc: 0.9508 - val_
Epoch 23/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1512 - acc: 0.9489 - val_
Epoch 24/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1434 - acc: 0.9513 - val_
Epoch 25/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1805 - acc: 0.9414 - val_
Epoch 26/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1453 - acc: 0.9528 - val_
Epoch 27/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1385 - acc: 0.9520 - val_
Epoch 28/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1420 - acc: 0.9533 - val_
Epoch 29/30
7352/7352 [=====] - 82s 11ms/step - loss: 0.1288 - acc: 0.9547 - val_

Epoch 30/30

7352/7352 [=====] - 82s 11ms/step - loss: 0.1291 - acc: 0.9532 - val_

```
In [26]: # Final evaluation of the model
scores4 = model4.evaluate(X_test, Y_test, verbose=0)
print("Test Score: %f" % (scores4[0]))
print("Test Accuracy: %f%%" % (scores4[1]*100))


# Confusion Matrix
Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])
Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model4.predict(X_test), axis=1)])

# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fontsize=12)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right', fontsize=12)
plt.ylabel('True label', size=18)
plt.xlabel('Predicted label', size=18)
plt.title("Confusion Matrix\n", size=24)
plt.show()
```

Test Score: 0.545492

Test Accuracy: 89.921955%



6 (6) Model having 2 LSTM layer with 64 LSTM Units and 'rmsprop' as an optimizer

```
In [27]: # Initiliazing the sequential model
model5 = Sequential()
# Configuring the parameters
model5.add(LSTM(64,return_sequences=True, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model5.add(Dropout(0.7))

# Configuring the parameters
model5.add(LSTM(64))
# Adding a dropout layer
model5.add(Dropout(0.7))
# Adding a dense output layer with sigmoid activation
model5.add(Dense(n_classes, activation='sigmoid'))
```

```

print(model5.summary())

# Compiling the model
model5.compile(loss='categorical_crossentropy',optimizer='rmsprop',metrics=['accuracy'])

# Training the model
history5 = model5.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
lstm_7 (LSTM)	(None, 128, 64)	18944
dropout_7 (Dropout)	(None, 128, 64)	0
lstm_8 (LSTM)	(None, 64)	33024
dropout_8 (Dropout)	(None, 64)	0
dense_6 (Dense)	(None, 6)	390

Total params: 52,358
 Trainable params: 52,358
 Non-trainable params: 0

None

Train on 7352 samples, validate on 2947 samples

Epoch 1/30
 7352/7352 [=====] - 110s 15ms/step - loss: 1.1611 - acc: 0.4890 - val_loss: 1.1611 - val_acc: 0.4890

Epoch 2/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.8021 - acc: 0.6549 - val_loss: 0.8021 - val_acc: 0.6549

Epoch 3/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.7392 - acc: 0.6670 - val_loss: 0.7392 - val_acc: 0.6670

Epoch 4/30
 7352/7352 [=====] - 109s 15ms/step - loss: 0.6489 - acc: 0.7338 - val_loss: 0.6489 - val_acc: 0.7338

Epoch 5/30
 7352/7352 [=====] - 107s 15ms/step - loss: 0.5545 - acc: 0.7625 - val_loss: 0.5545 - val_acc: 0.7625

Epoch 6/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.4380 - acc: 0.8164 - val_loss: 0.4380 - val_acc: 0.8164

Epoch 7/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.3323 - acc: 0.8966 - val_loss: 0.3323 - val_acc: 0.8966

Epoch 8/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.2380 - acc: 0.9301 - val_loss: 0.2380 - val_acc: 0.9301

Epoch 9/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.2048 - acc: 0.9370 - val_loss: 0.2048 - val_acc: 0.9370

Epoch 10/30
 7352/7352 [=====] - 108s 15ms/step - loss: 0.1991 - acc: 0.9389 - val_loss: 0.1991 - val_acc: 0.9389

Epoch 11/30

```

7352/7352 [=====] - 108s 15ms/step - loss: 0.1775 - acc: 0.9429 - val_
Epoch 12/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1724 - acc: 0.9436 - val_
Epoch 13/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1628 - acc: 0.9453 - val_
Epoch 14/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1754 - acc: 0.9440 - val_
Epoch 15/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1556 - acc: 0.9465 - val_
Epoch 16/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1733 - acc: 0.9452 - val_
Epoch 17/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1737 - acc: 0.9448 - val_
Epoch 18/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1698 - acc: 0.9442 - val_
Epoch 19/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1469 - acc: 0.9512 - val_
Epoch 20/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1492 - acc: 0.9461 - val_
Epoch 21/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1532 - acc: 0.9490 - val_
Epoch 22/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1758 - acc: 0.9436 - val_
Epoch 23/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1528 - acc: 0.9489 - val_
Epoch 24/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1589 - acc: 0.9438 - val_
Epoch 25/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1522 - acc: 0.9437 - val_
Epoch 26/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1363 - acc: 0.9495 - val_
Epoch 27/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1480 - acc: 0.9459 - val_
Epoch 28/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1347 - acc: 0.9495 - val_
Epoch 29/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1489 - acc: 0.9474 - val_
Epoch 30/30
7352/7352 [=====] - 108s 15ms/step - loss: 0.1491 - acc: 0.9486 - val_

```

```

In [28]: # Final evaluation of the model
         scores5 = model5.evaluate(X_test, Y_test, verbose=0)
         print("Test Score: %f" % (scores5[0]))
         print("Test Accuracy: %f%%" % (scores5[1]*100))

         # Confusion Matrix
         Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_test, axis=1)])

```

```

Y_predictions = pd.Series([ACTIVITIES[y] for y in np.argmax(model5.predict(X_test), axis=1)])


# Code for drawing seaborn heatmaps
class_names = ['laying', 'sitting', 'standing', 'walking', 'walking_downstairs', 'walking_upstairs']
df_heatmap = pd.DataFrame(confusion_matrix(Y_true, Y_predictions), index=class_names, columns=class_names)
fig = plt.figure(figsize=(10,7))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fontsize=12)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=90, ha='right', fontsize=12)
plt.ylabel('True label', size=18)
plt.xlabel('Predicted label', size=18)
plt.title("Confusion Matrix\n", size=24)
plt.show()

```

Test Score: 0.412691

Test Accuracy: 90.939939%



7 CONCLUSION

8 (a). Procedure Followed :

STEP 1 :- Load the data and split into training_data and test_data

STEP 2:-Try out different LSTM architectures

STEP 3:- Find test score and accuracy for each model

STEP 4:- Draw confusion matrix using seaborn heatmap for each model

9 (b). Table (Model performances) :

```
In [29]: # Creating table using PrettyTable library
         from prettytable import PrettyTable

         # Names of models
         names = ['1 LSTM layer with 32 LSTM Units(Optimizer-->rmsprop)', '1 LSTM layer with 48 LSTM Units(Optimizer-->rmsprop)', '1 LSTM layer with 64 LSTM Units(Optimizer-->rmsprop)', '1 LSTM layer with 32 LSTM Units(Optimizer-->rmsprop)', '1 LSTM layer with 48 LSTM Units(Optimizer-->rmsprop)', '1 LSTM layer with 64 LSTM Units(Optimizer-->rmsprop)']

         # Training accuracies
         train_acc = [history.history['acc'][29], history1.history['acc'][29], history2.history['acc'][29], history3.history['acc'][29], history4.history['acc'][29], history5.history['acc'][29]]

         # Test accuracies
         test_acc = [scores[1], scores1[1], scores2[1], scores3[1], scores4[1], scores5[1]]

         numbering = [1,2,3,4,5,6]

         # Initializing prettytable
         ptable = PrettyTable()

         # Adding columns
         ptable.add_column("S.NO.", numbering)
         ptable.add_column("MODEL", names)
         ptable.add_column("Training Accuracy", train_acc)
         ptable.add_column("Test Accuracy", test_acc)

         # Printing the Table
         print(ptable)
```

S.NO.	MODEL	Training Accuracy	Test Accuracy
1	1 LSTM layer with 32 LSTM Units(Optimizer-->rmsprop)	0.9411044613710555	0.882931

	2		1 LSTM layer with 48 LSTM Units(Optimizer-->adam)		0.9272306855277476		0.871394
	3		1 LSTM layer with 48 LSTM Units(Optimizer-->rmsprop)		0.9470892274211099		0.904988
	4		1 LSTM layer with 64 LSTM Units(Optimizer-->rmsprop)		0.948993471164309		0.920597
	5		2 LSTM layer with 32 LSTM Units(Optimizer-->rmsprop)		0.9532100108813928		0.899219
	6		2 LSTM layer with 64 LSTM Units(Optimizer-->rmsprop)		0.9485854189336235		0.909399
+-----+-----+-----+-----+-----+-----+-----+-----+							